

Charlie and the Reasoning Room

*Reason — Reinforcement Learning,
from bandits to reasoning models*




Fig. 1. From bandits to reasoning. *a* the arms of a bandit; *b* each arm's distribution of reward; *c* the arm pulled; *d* a tree of reasoning steps grown from it; *e* the path that reward reinforced.