

Charlie and the Reasoning Room

*Reason — Reinforcement Learning,
from bandits to reasoning models*




Fig. 1. From bandits to reasoning. *a* a bandit's arms; *b* each arm's reward on one scale, its mean ticked; *c* the arm with the highest mean, pulled; *d* a tree of reasoning grown from it; *e* the path reward reinforced.