

5D Parallelism for Large Model Training

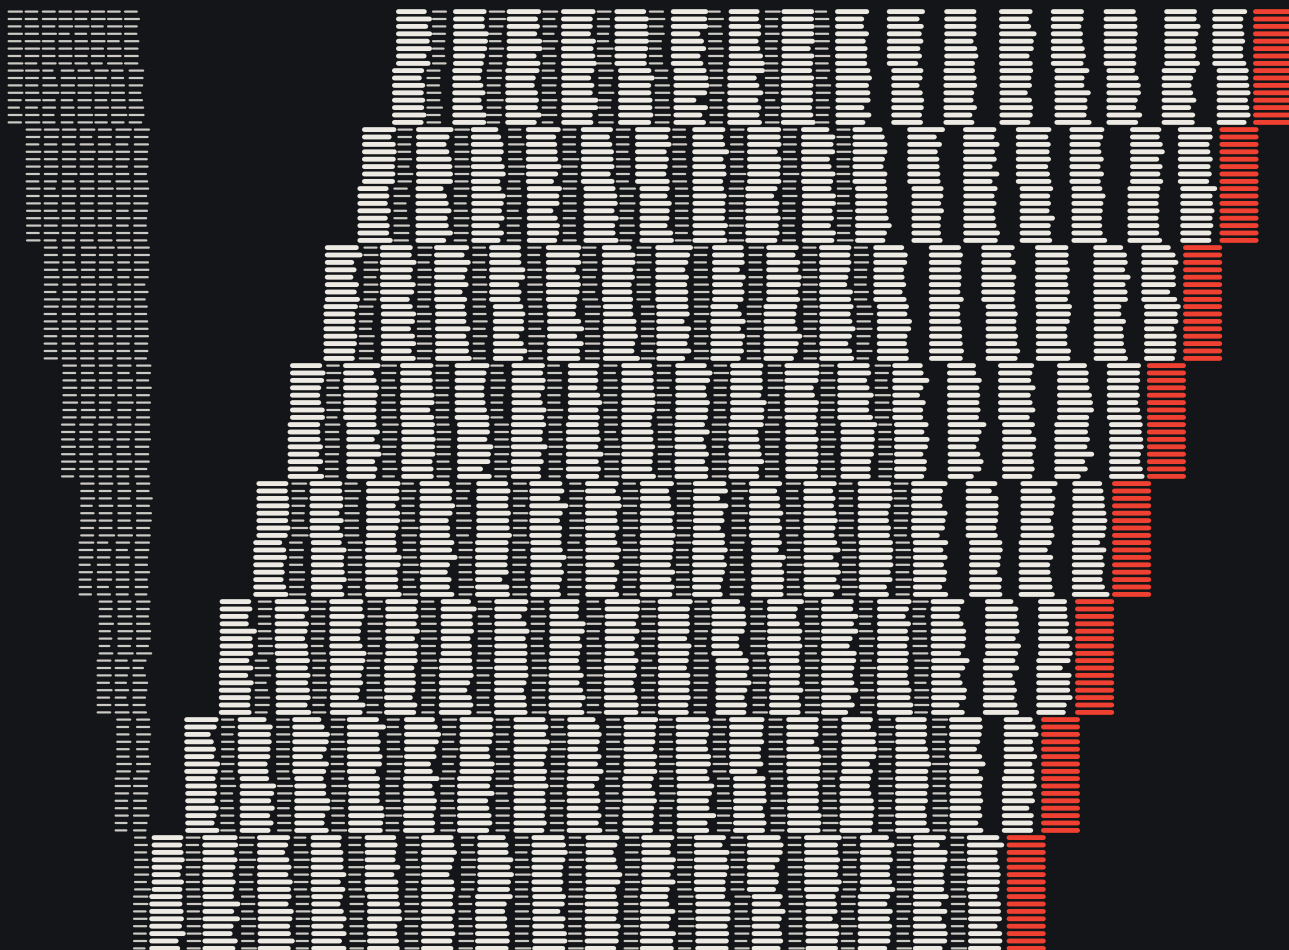


FIG. 0 One training step on a 128-GPU mesh – data 2 × pipeline 8 × tensor 2 × context 2 × expert 2 – simulated under a 1F1B schedule of 16 micro-batches with ±7% kernel-time noise. Each row is a GPU: thin = forward, thick = backward, red = gradient all-reduce.