

Build Large Language Models (LLMs) from Scratch




FIG. 0 The titles of all 52 Vizudara books, tokenised by a byte-pair encoder trained on them in this page: 195 merges turn 1,593 characters into 526 tokens. Each arc joins two uses of one token – tokens that begin a word above the line, the rest below; white is “Scratch”, used 11 times.