

Build Large Language Models (LLMs) from Scratch




FIG. 0 A byte-pair tokenizer trained in this page on the titles of all 52 Vizudara books (195 merges), then run on this one. Each arch is a merge, as tall as it was late to be learned; the 47 characters become the 10 tokens barred below. The merges that build "Scratch" are white.