

Charlie and the Vision Room

Vision – Vision Transformers




FIG. 0 A Vision Transformer reads a picture as a 12×12 grid of patches and tells them apart by position embeddings. These are its fixed 2-D sine-cosine ones ($d = 256$): each small field contours one patch's cosine similarity to every other position; white marks the patch itself.