

Charlie and the Vision Room

Vision – Vision Transformers




FIG. 0 A Vision Transformer reads a picture as an 8×8 grid of patches and tells them apart by position embeddings – here the fixed 2-D sine-cosine kind used by MAE and Simple ViT ($d = 256$). Each small field contours one patch's cosine similarity to every position; white marks the patch itself.