

FORTHCOMING

# Kernel Engineering

From silicon to speculative decoding —  
GPU kernels for modern LLMs.

