

FORTHCOMING

Kernel Engineering

From silicon to speculative decoding —
GPU kernels for modern LLMs.

