

```
[412] rate limits apply per workspace · [log] 14:02:11 cache miss ratio 0.31 · [doc] migration guide v3 to v4 · [ti
[et] p95 latency regression after deploy · [wiki] onboarding ·
[log] 14:02:13 upstream timeout 3000ms · [doc] retry policy: expon
webhook signatures · [log] GET /v1/orders 200 212ms · [doc]
pplly per workspace · [log] 14:02:11 cache miss ratio 0.31 · [d
sion after deploy · [wiki] onboarding checklist · [slack] can
eout 3
[log] 14
[ki] on
ry pol
212ms
ratio
t · [s
backof
iction
ration
review
Q3 in
e · [
[ti]
[log]
0977]
its ap
egress
m time
· [lo
· [wik
[retr
s 200
miss
cklist
tial b
is evi
[mi]
meone
memo]
plaine
v4 ·
PR? ·
[kb/0
e limi
ney re
stream
atures
space
play ·
[doc]
orders
cache
```

CONTEXT WINDOW · 128,000 TOKENS · 127,406 USED

SYSTEM

You are the on-call assistant for the billing platform. Answer only from the documents and tool results provided in this context. Cite every claim as [doc:n] or [tool:n]. If the context does not contain the answer, say exactly what is missing instead of guessing. Prefer short answers; show commands in code blocks; never run a destructive command without asking first; escalate to a human when a fix touches customer money. Timezone for all logs: UTC.

TOOLS

"name": "search\_logs",  
"description": "Full-text search over service logs", "input\_schema":  
{"type": "object", "properties":  
{"query": {"type": "string"},  
"since": {"type": "string"},  
"limit": {"type": "integer"},  
"required": [{"query"}]} {"name":  
"get\_metric", "description": "Fetch a time series", "input\_schema":  
{"type": "object", "properties":  
{"name": {"type": "string"},  
"window": {"type": "string"},  
"required": [{"name"}]}

DOCS

[doc:1] The response cache sits in front of the orders service; entries expire after 300 s. [doc:2] Cache keys are built from the tenant id and the request path; query strings are ignored unless listed in vary\_on. [doc:3] On a miss the service calls the upstream ledger, p95 180 ms, which is rate limited to 2,000 requests per second per region. [doc:5] Config changes to vary\_on take effect on the next deploy.

MEMORY

- user is on the payments team - prefers commands over prose - last week: rolled back cache v2 after key collisions - staging mirrors prod config

HISTORY

user: latency alarms again on orders. assistant: checking the dashboards [tool:1] p95 went from 190 ms to 410 ms at 13:58, right after deploy 8812. user: that deploy only touched the cache config. assistant: then misses are the first suspect; pulling the miss ratio [tool:2] 0.31, up from 0.04. user: so the cache is barely working? assistant: it is working, but on the wrong keys: [tool:3] deploy 8812 added the session cookie to vary\_on, so every user now gets a private copy of every response. user: which endpoints are affected? assistant: all GET routes under /v1/orders; the ledger is absorbing the misses, which is why p95 doubled rather than failing outright. user: can we just drop the cookie from vary\_on? assistant: yes for anonymous routes; two authenticated routes render per-user totals and need the user id in the key instead. user: ok, what else? assistant: purge the cache after the fix so the private copies do not linger for the full 300 s.

USER

why did p95 double after the cache change, and what is the smallest safe fix?

594 tokens free

```
ctx = pack(system, tools, docs, memory, history, user, budget=128_000)
```

# \$ AI Context Engineering ■

VIZUARA BOOKS ■

\$ INTERMEDIATE · 43 CAPSULES · 10 H